


```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import datetime
from sklearn.preprocessing import LabelEncoder
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import seaborn as sns
from keras.layers import Dense, BatchNormalization, Dropout, LSTM
from keras.models import Sequential
from keras.utils import to_categorical
from keras.optimizers import Adam
from tensorflow.keras import regularizers
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report, accuracy_score, f1_score
from keras import callbacks
```

```
df=pd.read_csv("weatherAUS.csv")
```

df

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGust
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	
...	
145455	2017-06-21	Uluru	2.8	23.4	0.0	NaN	NaN	
145456	2017-06-22	Uluru	3.6	25.3	0.0	NaN	NaN	
145457	2017-06-23	Uluru	5.4	26.9	0.0	NaN	NaN	
145458	2017-06-24	Uluru	7.8	27.0	0.0	NaN	NaN	
145459	2017-06-25	Uluru	14.9	NaN	0.0	NaN	NaN	

145460 rows × 23 columns



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   Date                 145460 non-null object
1   Location             145460 non-null object
2   MinTemp              143975 non-null float64
3   MaxTemp              144199 non-null float64
4   Rainfall             142199 non-null float64
5   Evaporation          82670 non-null float64
6   Sunshine             75625 non-null float64
7   WindGustDir          135134 non-null object
8   WindGustSpeed        135197 non-null float64
9   WindDir9am           134894 non-null object
10  WindDir3pm           141232 non-null object
11  WindSpeed9am         143693 non-null float64
12  WindSpeed3pm         142398 non-null float64
13  Humidity9am          142806 non-null float64
14  Humidity3pm          140953 non-null float64
15  Pressure9am          130395 non-null float64
16  Pressure3pm          130432 non-null float64
```

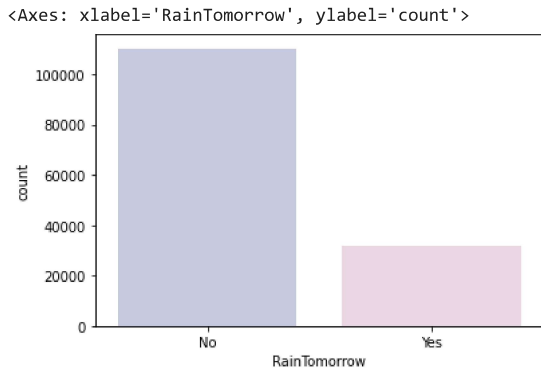
```
17 Cloud9am      89572 non-null float64
18 Cloud3pm      86102 non-null float64
19 Temp9am       143693 non-null float64
20 Temp3pm       141851 non-null float64
21 RainToday     142199 non-null object
22 RainTomorrow  142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB
```

```
#Points to notice:
```

```
#There are missing values in the dataset
#Dataset includes numeric and categorical values
```


```
#DATA VISUALIZATION AND CLEANING
```

```
cols= ["#C2C4E2", "#EED4E5"]
sns.countplot(x= df["RainTomorrow"], palette= cols)
```



```
corrmat = df.corr()
cmap = sns.diverging_palette(260,-10,s=50, l=75, n=6, as_cmap=True)
plt.subplots(figsize=(18,18))
sns.heatmap(corrmat, cmap= cmap, annot=True, square=True)
```

<Axes: >



```
lengths = df["Date"].str.len()
lengths.value_counts()
```

```
10      145460
Name: Date, dtype: int64
```

```
df['Date']= pd.to_datetime(df["Date"])
df['year'] = df.Date.dt.year
```

```
def encode(df, col, max_val):
    df[col + '_sin'] = np.sin(2 * np.pi * df[col]/max_val)
    df[col + '_cos'] = np.cos(2 * np.pi * df[col]/max_val)
    return df
```

```
df['month'] = df.Date.dt.month
df = encode(df, 'month', 12)
```

```
df['day'] = df.Date.dt.day
df = encode(df, 'day', 31)
```


```
df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindDir9am	...	Temp3pm	RainToday
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	W	...	21.8	N
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	NNW	...	24.3	N
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	W	...	23.2	N
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	SE	...	26.5	N
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	ENE	...	29.7	N

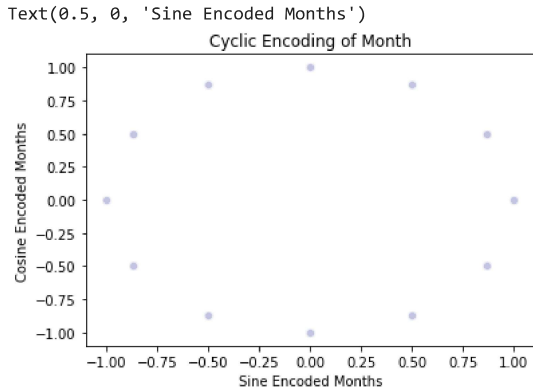
5 rows × 30 columns



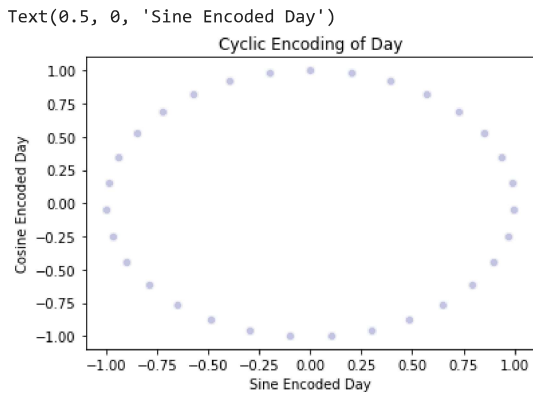
```
section = df[:360]
tm = section["day"].plot(color="#C2C4E2")
tm.set_title("Distribution Of Days Over Year")
tm.set_ylabel("Days In month")
tm.set_xlabel("Days In Year")
```



```
cyclic_month = sns.scatterplot(x="month_sin",y="month_cos",data=df, color="#C2C4E2")
cyclic_month.set_title("Cyclic Encoding of Month")
cyclic_month.set_ylabel("Cosine Encoded Months")
cyclic_month.set_xlabel("Sine Encoded Months")
```



```
cyclic_day = sns.scatterplot(x='day_sin',y='day_cos',data=df, color="#C2C4E2")
cyclic_day.set_title("Cyclic Encoding of Day")
cyclic_day.set_ylabel("Cosine Encoded Day")
cyclic_day.set_xlabel("Sine Encoded Day")
```



```
#Filling missing values with mode of the column value
```

```
#Get list of categorical variables
```

```
s = (df.dtypes == "object")
object_cols = list(s[s].index)
```

```
print("Categorical variables:")
print(object_cols)
```

```
Categorical variables:
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']
```

```
# Missing values in categorical variables
```

```
for i in object_cols:
    print(i, df[i].isnull().sum())
```

```
Location 0
WindGustDir 10326
WindDir9am 10566
WindDir3pm 4228
RainToday 3261
RainTomorrow 3267
```

```
# Filling missing values with mode of the column in value
```

```
for i in object_cols:
    df[i].fillna(df[i].mode()[0], inplace=True)
```

```
#Numerical variables
```

```
#Filling missing values with median of the column value
```

```
# Get list of neumeric variables
t = (df.dtypes == "float64")
```

```

num_cols = list(t[t].index)

print("Neumeric variables:")
print(num_cols)

Numeric variables:
['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humi

# Missing values in numeric variables

for i in num_cols:
    print(i, df[i].isnull().sum())

MinTemp 1485
MaxTemp 1261
Rainfall 3261
Evaporation 62790
Sunshine 69835
WindGustSpeed 10263
WindSpeed9am 1767
WindSpeed3pm 3062
Humidity9am 2654
Humidity3pm 4507
Pressure9am 15065
Pressure3pm 15028
Cloud9am 55888
Cloud3pm 59358
Temp9am 1767
Temp3pm 3609
month_sin 0
month_cos 0
day_sin 0
day_cos 0

# Filling missing values with median of the column in value

for i in num_cols:
    df[i].fillna(df[i].median(), inplace=True)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date                  145460 non-null  datetime64[ns]
 1   Location              145460 non-null  object
 2   MinTemp               145460 non-null  float64
 3   MaxTemp               145460 non-null  float64
 4   Rainfall              145460 non-null  float64
 5   Evaporation           145460 non-null  float64
 6   Sunshine              145460 non-null  float64
 7   WindGustDir           145460 non-null  object
 8   WindGustSpeed         145460 non-null  float64
 9   WindDir9am            145460 non-null  object
10   WindDir3pm            145460 non-null  object
11   WindSpeed9am          145460 non-null  float64
12   WindSpeed3pm          145460 non-null  float64
13   Humidity9am           145460 non-null  float64
14   Humidity3pm           145460 non-null  float64
15   Pressure9am           145460 non-null  float64
16   Pressure3pm           145460 non-null  float64
17   Cloud9am              145460 non-null  float64
18   Cloud3pm              145460 non-null  float64
19   Temp9am               145460 non-null  float64
20   Temp3pm               145460 non-null  float64
21   RainToday             145460 non-null  object
22   RainTomorrow          145460 non-null  object
23   year                  145460 non-null  int64
24   month                 145460 non-null  int64
25   month_sin             145460 non-null  float64
26   month_cos             145460 non-null  float64
27   day                   145460 non-null  int64
28   day_sin               145460 non-null  float64
29   day_cos               145460 non-null  float64
dtypes: datetime64[ns](1), float64(20), int64(3), object(6)
memory usage: 33.3+ MB

#DATA PREPROCESSING

#Label encoding the catagorical variable

# Apply label encoder to each column with categorical data

```

```

label_encoder = LabelEncoder()
for i in object_cols:
    df[i] = label_encoder.fit_transform(df[i])

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   145460 non-null  datetime64[ns]
1   Location               145460 non-null  int64
2   MinTemp                145460 non-null  float64
3   MaxTemp                145460 non-null  float64
4   Rainfall               145460 non-null  float64
5   Evaporation            145460 non-null  float64
6   Sunshine               145460 non-null  float64
7   WindGustDir            145460 non-null  int64
8   WindGustSpeed          145460 non-null  float64
9   WindDir9am             145460 non-null  int64
10  WindDir3pm             145460 non-null  int64
11  WindSpeed9am           145460 non-null  float64
12  WindSpeed3pm           145460 non-null  float64
13  Humidity9am            145460 non-null  float64
14  Humidity3pm            145460 non-null  float64
15  Pressure9am            145460 non-null  float64
16  Pressure3pm            145460 non-null  float64
17  Cloud9am                145460 non-null  float64
18  Cloud3pm               145460 non-null  float64
19  Temp9am                145460 non-null  float64
20  Temp3pm                145460 non-null  float64
21  RainToday              145460 non-null  int64
22  RainTomorrow           145460 non-null  int64
23  year                   145460 non-null  int64
24  month                  145460 non-null  int64
25  month_sin              145460 non-null  float64
26  month_cos              145460 non-null  float64
27  day                    145460 non-null  int64
28  day_sin                145460 non-null  float64
29  day_cos                145460 non-null  float64
dtypes: datetime64[ns](1), float64(20), int64(9)
memory usage: 33.3 MB

```

```
# Preparing attributes of scale data
```

```
features = df.drop(['RainTomorrow', 'Date', 'day', 'month'], axis=1) # dropping target and extra columns
```

```
target = df['RainTomorrow']
```

```
#Set up a standard scaler for the features
```

```
col_names = list(features.columns)
```

```
s_scaler = preprocessing.StandardScaler()
```

```
features = s_scaler.fit_transform(features)
```


```
features = pd.DataFrame(features, columns=col_names)
```

```
features.describe().T
```

	count	mean	std	min	25%	50%	75%	max
Location	145460.0	7.815677e-18	1.000003	-1.672228	-0.899139	0.014511	0.857881	1.701250
MinTemp	145460.0	-4.501830e-16	1.000003	-3.250525	-0.705659	-0.030170	0.723865	3.410112
MaxTemp	145460.0	3.001220e-16	1.000003	-3.952405	-0.735852	-0.086898	0.703133	3.510563
Rainfall	145460.0	7.815677e-18	1.000003	-0.275097	-0.275097	-0.275097	-0.203581	43.945571
Evaporation	145460.0	-3.282584e-17	1.000003	-1.629472	-0.371139	-0.119472	0.006361	43.985108
Sunshine	145460.0	-5.424080e-16	1.000003	-2.897217	0.076188	0.148710	0.257494	2.360634
WindGustDir	145460.0	6.252542e-18	1.000003	-1.724209	-0.872075	0.193094	1.045228	1.471296
WindGustSpeed	145460.0	1.824961e-16	1.000003	-2.588407	-0.683048	-0.073333	0.460168	7.243246
WindDir9am	145460.0	7.190423e-17	1.000003	-1.550000	-0.885669	0.000105	0.885879	1.771653
WindDir3pm	145460.0	8.284618e-17	1.000003	-1.718521	-0.837098	0.044324	0.925747	1.586813
WindSpeed9am	145460.0	5.627087e-17	1.000003	-1.582204	-0.702280	0.112214	0.560750	13.086470



```
#Detecting outliers
#looking at the scaled features
colours = ["#D0DBEE", "#C2C4E2", "#EED4E5", "#D1E6DC", "#BDE2E2"]
plt.figure(figsize=(20,10))
sns.boxenplot(data= features,palette = colours)
plt.xticks(rotation=90)
plt.show()
```



```
features["RainTomorrow"] = target

X = features.drop(["RainTomorrow"], axis=1)
y = features["RainTomorrow"]

# Splitting test and training sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

X.shape

(145460, 26)

#Early stopping
early_stopping = callbacks.EarlyStopping(
    min_delta=0.001, # minimum amount of change to count as an improvement
    patience=20, # how many epochs to wait before stopping
    restore_best_weights=True,
```

```

)

# Initialising the NN
model = Sequential()

# layers

model.add(Dense(units = 32,activation = 'relu', input_dim = 26))
model.add(Dense(units = 32,activation = 'relu'))
model.add(Dense(units = 16,activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 1,activation = 'sigmoid'))

# Compiling the ANN
opt = Adam(learning_rate=0.00009)
model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])

# Train the ANN
history = model.fit(X_train, y_train, batch_size = 32, epochs=50, callbacks=[early_stopping], validation_split=0.2)

Epoch 28/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3750 - accuracy: 0.8339 - val_loss: 0.3535 - val_accuracy: 0.8
Epoch 29/150
2910/2910 [=====] - 11s 4ms/step - loss: 0.3771 - accuracy: 0.8337 - val_loss: 0.3548 - val_accuracy: 0.
Epoch 30/150
2910/2910 [=====] - 8s 3ms/step - loss: 0.3742 - accuracy: 0.8330 - val_loss: 0.3549 - val_accuracy: 0.8
Epoch 31/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3745 - accuracy: 0.8332 - val_loss: 0.3533 - val_accuracy: 0.8
Epoch 32/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3768 - accuracy: 0.8330 - val_loss: 0.3533 - val_accuracy: 0.8
Epoch 33/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3755 - accuracy: 0.8328 - val_loss: 0.3528 - val_accuracy: 0.
Epoch 34/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3745 - accuracy: 0.8337 - val_loss: 0.3534 - val_accuracy: 0.8
Epoch 35/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3746 - accuracy: 0.8342 - val_loss: 0.3531 - val_accuracy: 0.8
Epoch 36/150
2910/2910 [=====] - 11s 4ms/step - loss: 0.3742 - accuracy: 0.8331 - val_loss: 0.3524 - val_accuracy: 0.
Epoch 37/150
2910/2910 [=====] - 10s 4ms/step - loss: 0.3737 - accuracy: 0.8334 - val_loss: 0.3533 - val_accuracy: 0.
Epoch 38/150
2910/2910 [=====] - 8s 3ms/step - loss: 0.3735 - accuracy: 0.8341 - val_loss: 0.3530 - val_accuracy: 0.8
Epoch 39/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3740 - accuracy: 0.8332 - val_loss: 0.3543 - val_accuracy: 0.
Epoch 40/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3753 - accuracy: 0.8334 - val_loss: 0.3532 - val_accuracy: 0.8
Epoch 41/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3742 - accuracy: 0.8331 - val_loss: 0.3541 - val_accuracy: 0.
Epoch 42/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3741 - accuracy: 0.8339 - val_loss: 0.3535 - val_accuracy: 0.8
Epoch 43/150
2910/2910 [=====] - 12s 4ms/step - loss: 0.3725 - accuracy: 0.8342 - val_loss: 0.3524 - val_accuracy: 0.
Epoch 44/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3738 - accuracy: 0.8328 - val_loss: 0.3529 - val_accuracy: 0.8
Epoch 45/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3726 - accuracy: 0.8346 - val_loss: 0.3529 - val_accuracy: 0.
Epoch 46/150
2910/2910 [=====] - 8s 3ms/step - loss: 0.3722 - accuracy: 0.8334 - val_loss: 0.3539 - val_accuracy: 0.8
Epoch 47/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3719 - accuracy: 0.8339 - val_loss: 0.3526 - val_accuracy: 0.
Epoch 48/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3751 - accuracy: 0.8328 - val_loss: 0.3540 - val_accuracy: 0.
Epoch 49/150
2910/2910 [=====] - 8s 3ms/step - loss: 0.3719 - accuracy: 0.8333 - val_loss: 0.3538 - val_accuracy: 0.8
Epoch 50/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3733 - accuracy: 0.8342 - val_loss: 0.3525 - val_accuracy: 0.8
Epoch 51/150
2910/2910 [=====] - 11s 4ms/step - loss: 0.3725 - accuracy: 0.8339 - val_loss: 0.3529 - val_accuracy: 0.
Epoch 52/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3738 - accuracy: 0.8334 - val_loss: 0.3530 - val_accuracy: 0.8
Epoch 53/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3735 - accuracy: 0.8344 - val_loss: 0.3527 - val_accuracy: 0.8
Epoch 54/150
2910/2910 [=====] - 10s 3ms/step - loss: 0.3719 - accuracy: 0.8338 - val_loss: 0.3531 - val_accuracy: 0.
Epoch 55/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3731 - accuracy: 0.8347 - val_loss: 0.3526 - val_accuracy: 0.8
Epoch 56/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3722 - accuracy: 0.8342 - val_loss: 0.3539 - val_accuracy: 0.8
Epoch 57/150
2910/2910 [=====] - 9s 3ms/step - loss: 0.3722 - accuracy: 0.8342 - val_loss: 0.3539 - val_accuracy: 0.8

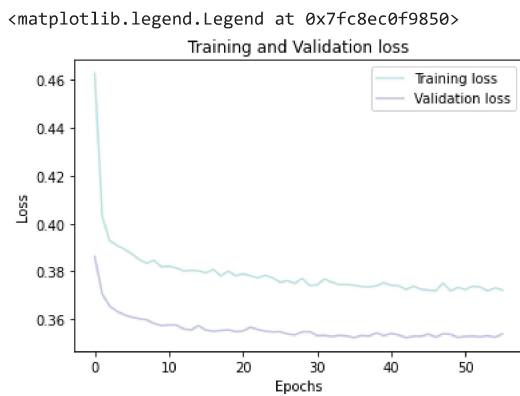
history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['loss']], "#BDE2E2", label='Training loss')
plt.plot(history_df.loc[:, ['val_loss']], "#C2C4E2", label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')

```




```
plt.legend(loc="best")
```



```
history_df = pd.DataFrame(history.history)
```

```
plt.plot(history_df.loc[:, ['accuracy']], "#BDE2E2", label='Training accuracy')
plt.plot(history_df.loc[:, ['val_accuracy']], "#C2C4E2", label='Validation accuracy')
```

```
plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
# Predicting the test set results
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)
# confusion matrix
cmap1 = sns.diverging_palette(260,-10,s=50, l=75, n=5, as_cmap=True)
plt.subplots(figsize=(12,8))
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix/np.sum(cf_matrix), cmap = cmap1, annot = True, annot_kws = {'size':15})
```

```
910/910 [=====] - 2s 2ms/step
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.86	0.96	0.91	22672
1	0.75	0.45	0.56	6420
accuracy			0.85	29092
macro avg	0.80	0.70	0.73	29092
weighted avg	0.84	0.85	0.83	29092

